

Geometric Transformation


Image Interpolation

- **Interpolation in image processing** is a technique used to estimate pixel values at non-sampled points based on known pixel values. It is essential in various applications, such as resizing images, enhancing image quality, and remapping pixel grids. Common methods of interpolation include:

- **Nearest Neighbour:** Assigns the value of the nearest known pixel to the unknown pixel.
- **Bilinear:** Applies linear interpolation in two dimensions, considering the closest four pixels.
- **Bicubic:** Uses the values of the nearest 16 pixels to calculate the new pixel value, providing smoother results.
- Interpolation is crucial in digital photos, especially during processes like Bayer demosaicing and photo enlargement. For a more detailed understanding, you can refer to the foundational concepts of image interpolation.
-

Nearest Neighbour

```
In [1]: from scipy.interpolate import interp1d  
In [2]: x = np.linspace(0, 6, 200); # fine sampling to represent continuous function  
In [3]: xs = np.array([0, 1, 2, 3, 4, 5, 6]); # the sample points  
In [4]: f = np.sin(x); F = np.sin(xs); # the 'continuous' function and its sampled version  
In [5]: ifunc = interp1d(xs, F, kind='nearest'); hatf_nn = ifunc(x);  
In [6]: plt.clf(); plt.plot(x, f, 'g-', xs, F, 'ro', x, hatf_nn, 'b-');
```




Linear Interpolation

Linear interpolation is the scientific equivalent of what you have already learned in kindergarten: connect the dots. Between two adjacent sample points k and $k + 1$ we assume the function is a linear function and thus in this interval $[k, k + 1]$ we have:


$$k \leq x \leq k + 1: \quad \hat{f}(x) = (1 - (x - k))F(k) + (x - k)F(k + 1)$$

```
In [7]: ifunc = interp1d(xs, F, kind='linear'); hatf_lin = ifunc(x);  
In [8]: plt.clf(); plt.plot(x, f, 'g-', xs, F, 'ro', x, hatf_lin, 'b-');
```




Cubic Interpolation

```
In [9]: ifunc = interp1d(xs, F, kind='cubic'); hatf_cubic = ifunc(x);  
In [10]: plt.clf(); plt.plot(x, f, 'g-', xs, F, 'ro', x, hatf_cubic, 'b-');
```



Geometric Operations

- **Scale** - change image content size
- **Rotate** - change image content orientation
- **Reflect** - flip over image contents
- **Translate** - change image content position
- **Affine Transformation**
 - general image content linear geometric transformation



Translation:



```
import numpy as np
import matplotlib.pyplot as plt

# Load the image
image = cv2.imread('D://SJC//Image Processing//images//image1.jpg') # Provide the path to your image here
image = cv2.cvtColor(image, cv2.COLOR_BGR2RGB)
cv2.imshow('Image', image) # 'Image' is the window name, and image is the image to be displayed
cv2.waitKey(0) # Wait indefinitely for a key press
cv2.destroyAllWindows()

# Translation
def translate_image(image, tx, ty):
    rows, cols, _ = image.shape
    # Define the translation matrix
    translation_matrix = np.float32([[1, 0, tx], [0, 1, ty]])
    # Apply the transformation using cv2.warpAffine
    translated_image = cv2.warpAffine(image, translation_matrix, (cols, rows))
    return translated_image

# Example: Translate the image 50 pixels right and 30 pixels down
translated_image = translate_image(image, 50, 30)
```

Translated Image




Rotation:

```
#rotation
def rotate_image(image, angle):
    rows, cols, _ = image.shape
    # Get the rotation matrix
    rotation_matrix = cv2.getRotationMatrix2D((cols / 2, rows / 2), angle, 1)
    # Apply the rotation
    rotated_image = cv2.warpAffine(image, rotation_matrix, (cols, rows))
    return rotated_image

# Example: Rotate the image by 45 degrees
rotated_image = rotate_image(image, 90)

# Display the result
plt.imshow(rotated_image)
plt.title('Rotated Image')
plt.axis('off')
plt.show()
```

Rotated Image



Scaling:

```
#scaling
```

```
def scale_image(image, fx, fy):
```

```
    # fx and fy are scaling factors along the x and y axes
```

```
    scaled_image = cv2.resize(image, None, fx=fx, fy=fy, interpolation=cv2.INTER_LINEAR)
```

```
    return scaled_image
```

```
# Example: Scale the image by a factor of 1.5 in both x and y directions
```

```
scaled_image = scale_image(image, 1.5, 1.5)
```

```
# Display the result
```

```
plt.imshow(scaled_image)
```

```
plt.title('Scaled Image')
```

```
plt.axis('off')
```

```
plt.show()
```


Scaled Image



Affine transform:

- An affine transformation is a type of geometric transformation which preserves collinearity (if a collection of points sits on a line before the transformation, they all sit on a line afterwards) and the ratios of distances between points on a line.

*# Affine Transformation Affine transformations preserve parallelism and ratios of distances.
#These can include translation, scaling, rotation, and shearing combined.*

```
def affine_transform(image):  
    rows, cols, _ = image.shape  
    # Define three points for the original image  
    pts1 = np.float32([[50, 50], [200, 50], [50, 200]])  
    # Define the corresponding points in the transformed image  
    pts2 = np.float32([[10, 100], [200, 50], [100, 250]])  
    # Get the affine transformation matrix  
    affine_matrix = cv2.getAffineTransform(pts1, pts2)  
    # Apply the affine transformation  
    affine_image = cv2.warpAffine(image, affine_matrix, (cols, rows))  
    return affine_image
```


Apply affine transformation
affine_image = affine_transform(image)

Display the result
plt.imshow(affine_image)
plt.title('Affine Transformed Image')
plt.axis('off')
plt.show()


Affine Transformed Image




Original Image



Translated Image




Rotated Image



Scaled Image




Affine Transformed Image




Geometric Transformations

- A geometric transform consists of two basic steps ...
 - Step1: determining the pixel co-ordinate transformation
 - mapping of the co-ordinates of the moving image pixel to the point in the fixed image.




Geometric transformations


- Step2: determining the brightness of the points in the digital grid of the transformed image.
 - brightness is usually computed as an interpolation of the brightnesses of several points in the neighborhood.



Fixed Image



Moving Image



xformed Moving Image

We'll discuss step 2 first.

Affine Transformation

- An affine transformation maps variables (e.g. pixel intensity values located at position in an input image) into new variables (e.g. in an output image) by applying a linear combination of translation, rotation, scaling operations.
- **Significance:** In some imaging systems, images are subject to geometric distortions. Applying an affine transformation to a uniformly distorted image can correct for a range of perspective distortions.



Affine Transformation (con'd)

- By defining only the B matrix, this transformation can carry out pure **translation**:


$$A = \begin{bmatrix} 1 & 0 \\ 0 & 1 \end{bmatrix}, B = \begin{bmatrix} b_1 \\ b_2 \end{bmatrix}$$

- Pure **rotation** uses the A matrix and is defined as (for positive angles being clockwise rotations):

$$A = \begin{bmatrix} \cos(\theta) & -\sin(\theta) \\ \sin(\theta) & \cos(\theta) \end{bmatrix}, B = \begin{bmatrix} 0 \\ 0 \end{bmatrix}$$

Affine Transformation (con'd)

- Pure **scaling** is defined as $A = \begin{bmatrix} a_{11} & 0 \\ 0 & a_{22} \end{bmatrix}, B = \begin{bmatrix} 0 \\ 0 \end{bmatrix}$
- Since the general affine transformation is defined by 6 constants, it is possible to define this transformation by specifying 3 corresponding point pairs (more in next class).



Affine Transformation (con'd)

- An affine transformation is equivalent to the composed effects of translation, rotation and scaling, and shearing.
- The general affine transformation is commonly expressed as below:


$$\begin{bmatrix} x_2 \\ y_2 \end{bmatrix} = A \begin{bmatrix} x_1 \\ y_1 \end{bmatrix} + B$$

0th order coefficients

1st order coefficients

Another Example

Interpolation on an image (4x4 \rightarrow 8x8) after scaling




Bilinear Interpolation

- Substituting with the values just obtained:

$$f(x', y') = \lambda(\mu f(x+1, y+1) + (1-\mu)f(x+1, y)) \\ + (1-\lambda)(\mu f(x, y+1) + (1-\mu)f(x, y))$$

- You can do the expansion as an exercise.
- This is the formulation for **bilinear interpolation**



Digital Image Processing

Lecture 6: Image Geometry

Prof. Charlene Tsai


Example

```
I = imread('cman.tif');  
tform = maketform('affine',[1 0 0; .5 1 0; 0 0 1]);  
J = imtransform(I,tform);  
imshow(I), figure, imshow(J)
```



General Interpolation: 0th and 1st orders

- Consider 2 functions $R_0(u)$ and $R_1(u)$



$$R_0(u) = \begin{cases} 0 & \text{if } u \leq -0.5 \\ 1 & \text{if } -0.5 < u \leq 0.5 \\ 0 & \text{if } u > 0.5 \end{cases}$$


$$R_1(u) = \begin{cases} 1+u & \text{if } u \leq 0 \\ 1-u & \text{if } u \geq 0 \end{cases}$$

Substitute $R_0(u)$ for $R(u)$ \Rightarrow nearest neighbour interpolation.

Substitute $R_1(u)$ for $R(u)$ \Rightarrow linear interpolation.

General Interpolation: 3rd order (Cubic)

$$R_3(u) = \begin{cases} 1.5|u|^3 - 2.5|u|^2 + 1 & \text{if } |u| \leq 1 \\ -0.5|u|^3 + 2.5|u|^2 - 4|u| + 2 & \text{if } 1 < |u| \leq 2 \end{cases}$$




$$f(x') = R_3(-1-\lambda)f(x_1) + R_3(-\lambda)f(x_2) + R_3(1-\lambda)f(x_3) + R_3(2-\lambda)f(x_4)$$

General Interpolation

- We wish to interpolate a value $f(x')$ for $x_1 \leq x' \leq x_2$ and suppose $x' - x_1 = \lambda$ ← $0 \leq \lambda \leq 1$


- We define an interpolated value $R(u)$ and set

$$f(x') = R(-\lambda)f(x_1) + R(1-\lambda)f(x_2)$$




General Interpolation: Bicubic

- **Step 2:** the fractional part of the calculated pixel's address in the y-direction is used to fit another cubic polynomial down the column, based on the interpolated brightness values that lie on the curves $F(i)$, $i = 0, \dots, 3$.




General Interpolation: Bicubic (2D)


- Bicubic interpolation fits a series of cubic polynomials to the brightness values contained in the 4 x 4 array of pixels surrounding the calculated address.
- **Step 1:** four cubic polynomials $F(i)$, $i = 0, 1, 2, 3$ are fit to the control points along the rows. The fractional part of the calculated pixel's address in the x-direction is used.



General Interpolation: Example



- Original detailed part of flower image (8bit, 75×75)
- Detailed part of super-resolution image (8bit, 300×300) :



NN Interpolation



Bilinear Interpolation



Bicubic Interpolation

```
im = imread('flower.jpg');  
im2= imresize(im,[800,800],method);
```


General Interpolation: Summary


- For NN interpolation, the output pixel is assigned the value of the pixel that the point falls within. No other pixels are considered.
- For bilinear interpolation, the output pixel value is a weighted average of pixels in the nearest 2-by-2 neighborhood.
- For bicubic interpolation, the output pixel value is a weighted average of pixels in the nearest 4-by-4 neighborhood.
- Bilinear method takes longer than nearest neighbor interpolation, and the bicubic method takes longer than bilinear.
- The greater the number of pixels considered, the more accurate the computation is, so there is a trade-off between processing time and quality.

Geometric transformations


- Geometric transformations are common in computer graphics, and are often used in image analysis.
- Geometric transforms permit the elimination of geometric distortion that occurs when an image is captured.
- If one attempts to match two different images of the same object, a geometric transformation may be needed.
- Examples?

Interpolation: Linear (1D)

■ General idea:




To calculate the interpolated values



$$\frac{F - f(x_1)}{\lambda} = \frac{f(x_2) - f(x_1)}{1}$$


Interpolation: Linear (2D)

- How a 4x4 image would be interpolated to produce an 8x8 image?



Interpolation: Nearest Neighbor

1-D




We assign $f(x'_i) = f(x_j)$

x_j is the original point closest to x'_i

The original function values

The interpolated values

2-D



Matlab *imtransform*

- The `imtransform` function accepts two primary arguments:
 - The image to be transformed
 - A spatial transformation structure, called a TFORM, that specifies the type of transformation you want to perform
- Specify the type of transformation in a TFORM structure.

Two ways to create a TFORM struct:



- Using the `maketform` function
- Using the `cp2tform` function

Rotation Operation


- A geometric transform which maps the position of a picture element in an input image onto a position in an output image by rotating it through an angle about an origin.
- Commonly used to improve the visual appearance of an image.
- Can also be useful as a pre-processor in applications where directional operators are involved.
- Rotation is a special case of affine transformation.

Rotation Operation: Example

```
I = imread('ic.tif');  
J = imrotate(I,35,'bilinear');  
imshow(I)  
figure, imshow(J)
```




Rotation Operation: Remedies (con'd)



(x', y') in rotated image

(x'', y'') is the rotated (x', y')
back into the original image


- The grey value at (x'', y'') can be found by interpolation.
- This value is the grey value for the pixel at (x', y') in the rotated image.




Rotation Operation (cont)

- Mapping of a point (x,y) to another (x',y') through a counter-clockwise rotation of θ

$$\begin{bmatrix} x' \\ y' \end{bmatrix} = \begin{pmatrix} \cos \theta & -\sin \theta \\ \sin \theta & \cos \theta \end{pmatrix} \begin{bmatrix} x \\ y \end{bmatrix}$$



$$\begin{bmatrix} x \\ y \end{bmatrix} = \begin{pmatrix} \cos \theta & \sin \theta \\ -\sin \theta & \cos \theta \end{pmatrix} \begin{bmatrix} x' \\ y' \end{bmatrix}$$



Scaling Operation

- To shrink or zoom the size of an image (or part of an image).
- To change the visual appearance of an image;
- To alter the quantity of information
- To use as a low-level pre-processor in multi-stage image processing chain which operates on features of a particular scale.
- Scaling is a special case of *affine transformation*.
- The matlab command is simply “imresize”.

Summary

- Interpolation of intensity values on non-grid points:
 - Nearest Neighbor (NN)
 - Bilinear
 - Bicubic
- Image transformation
 - Computation of intensity values of the transformed image
 - Discussed some instances of affine transformation
 - Translation
 - Rotation
 - Scaling

Using *maketform*

- When using the `maketform` function, you can specify the type of transformation, e.g.
 - 'affine'
 - 'projective'
 - 'composite', et al
 - 'custom' and 'composite' capabilities of `maketform` allow a virtually limitless variety of spatial transformations to be used
- Once you define the transformation in a `TFORM` struct, you can perform the transformation by calling `imtransform`.